Module 7: Sentiment and Emotion Analysis Lab

ITAI 2373 - Natural Language Processing

Lab Overview

Welcome to the Sentiment and Emotion Analysis lab! In this hands-on session, you'll build a complete emotion detection system that works with both text and speech. This lab connects directly to Module 7's concepts and gives you practical experience with real-world emotion analysis.

What You'll Build Today:

- 1. Text Sentiment Analyzer using VADER and TextBlob
- 2. Machine Learning Classifier with scikit-learn
- 3. Speech Emotion Detector using audio features
- 4. Multimodal System combining text and speech analysis

Real Data You'll Use:

- · Customer reviews from multiple domains
- Social media posts with emotion labels
- · Audio recordings with emotional speech
- · Multimodal datasets combining text and audio

Learning Objectives:

By the end of this lab, you will:

- · Understand the differences between rule-based and ML approaches to sentiment analysis
- · Build and evaluate multiple sentiment analysis systems
- · Extract and analyze emotional features from speech
- Create a multimodal emotion detection system
- · Critically evaluate bias and fairness in emotion detection systems

Before we start let's see how this lab ties to previous modules

Building on Previous Modules

Welcome to Module 7! This lab builds directly on everything you've learned so far:

From Module 1 (Introduction to NLP):

- We'll apply NLP to understand human emotions and opinions
- Real-world applications: customer feedback, social media monitoring, healthcare

From Module 2 (Text Preprocessing):

- · We'll use tokenization, normalization, and cleaning techniques
- · Preprocessing becomes crucial for emotion detection accuracy

From Module 3 (Audio and Preprocessing):

- · We'll extract emotional features from speech signals
- Combine text and audio for multimodal emotion analysis

From Module 4 (Text Representation):

- We'll use TF-IDF vectorization for machine learning approaches
- · Compare rule-based vs. ML feature representations

From Module 5 (Part-of-Speech Tagging):

- POS tags help identify emotional intensity (adjectives, adverbs)
- · Grammatical patterns reveal sentiment structure

From Module 6 (Syntax and Parsing):

- Dependency relationships show emotional targets ("I love THIS product")
- Syntactic patterns help resolve sentiment scope and negation

Run each cell in order and complete the exercises marked with / YOUR TURN

Part 0: Setup and Installation

First, let's install all the libraries we'll need for this comprehensive lab. We'll be working with text processing, machine learning, and audio analysis libraries.

```
# Install required libraries for Google Colab
!pip install vaderSentiment textblob librosa soundfile
!python -m textblob.download_corpora
Requirement already satisfied: vaderSentiment in /usr/local/lib/python3.11/dist-packages (3.3.2)
     Requirement already satisfied: textblob in /usr/local/lib/python3.11/dist-packages (0.19.0)
     Requirement already satisfied: librosa in /usr/local/lib/python3.11/dist-packages (0.11.0)
     Requirement already satisfied: soundfile in /usr/local/lib/python3.11/dist-packages (0.13.1)
     Requirement already satisfied: requests in /usr/local/lib/python3.11/dist-packages (from vaderSentiment) (2.32.3)
     Requirement already satisfied: nltk>=3.9 in /usr/local/lib/python3.11/dist-packages (from textblob) (3.9.1)
     Requirement already satisfied: audioread>=2.1.9 in /usr/local/lib/python3.11/dist-packages (from librosa) (3.0.1)
     Requirement already satisfied: numba>=0.51.0 in /usr/local/lib/python3.11/dist-packages (from librosa) (0.60.0)
     Requirement already satisfied: numpy>=1.22.3 in /usr/local/lib/python3.11/dist-packages (from librosa) (2.0.2)
     Requirement already satisfied: scipy>=1.6.0 in /usr/local/lib/python3.11/dist-packages (from librosa) (1.15.3)
     Requirement already satisfied: scikit-learn>=1.1.0 in /usr/local/lib/python3.11/dist-packages (from librosa) (1.6.1)
     Requirement already satisfied: joblib>=1.0 in /usr/local/lib/python3.11/dist-packages (from librosa) (1.5.1)
     Requirement already satisfied: decorator>=4.3.0 in /usr/local/lib/python3.11/dist-packages (from librosa) (4.4.2)
     Requirement already satisfied: pooch>=1.1 in /usr/local/lib/python3.11/dist-packages (from librosa) (1.8.2)
     Requirement already satisfied: soxr>=0.3.2 in /usr/local/lib/python3.11/dist-packages (from librosa) (0.5.0.post1)
     Requirement already satisfied: typing extensions>=4.1.1 in /usr/local/lib/python3.11/dist-packages (from librosa) (4.14.1)
     Requirement already satisfied: lazy loader>=0.1 in /usr/local/lib/python3.11/dist-packages (from librosa) (0.4)
     Requirement already satisfied: msgpack>=1.0 in /usr/local/lib/python3.11/dist-packages (from librosa) (1.1.1)
     Requirement already satisfied: cffi>=1.0 in /usr/local/lib/python3.11/dist-packages (from soundfile) (1.17.1)
     Requirement already satisfied: pycparser in /usr/local/lib/python3.11/dist-packages (from cffi>=1.0->soundfile) (2.22)
     Requirement already satisfied: packaging in /usr/local/lib/python3.11/dist-packages (from lazy loader>=0.1->librosa) (24.2)
     Requirement already satisfied: click in /usr/local/lib/python3.11/dist-packages (from nltk>=3.9->textblob) (8.2.1)
     Requirement already satisfied: regex>=2021.8.3 in /usr/local/lib/python3.11/dist-packages (from nltk>=3.9->textblob) (2024.11.6)
     Requirement already satisfied: tqdm in /usr/local/lib/python3.11/dist-packages (from nltk>=3.9->textblob) (4.67.1)
     Requirement already satisfied: llvmlite<0.44,>=0.43.0dev0 in /usr/local/lib/python3.11/dist-packages (from numba>=0.51.0->librosa) (0.43.0)
     Requirement already satisfied: platformdirs>=2.5.0 in /usr/local/lib/python3.11/dist-packages (from pooch>=1.1->librosa) (4.3.8)
     Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/dist-packages (from requests->vaderSentiment) (3.4.2)
     Requirement already satisfied: idna<4.>=2.5 in /usr/local/lib/python3.11/dist-packages (from requests->yaderSentiment) (3.10)
     Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11/dist-packages (from requests->vaderSentiment) (2.4.0)
     Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.11/dist-packages (from requests->vaderSentiment) (2025.6.15)
     Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.11/dist-packages (from scikit-learn>=1.1.0->librosa) (3.6.0)
```

```
[nltk data] Downloading package brown to /root/nltk data...
     [nltk_data] Package brown is already up-to-date!
     [nltk data] Downloading package punkt tab to /root/nltk data...
     [nltk_data] Package punkt_tab is already up-to-date!
     [nltk_data] Downloading package wordnet to /root/nltk_data...
     [nltk data] Package wordnet is already up-to-date!
     [nltk data] Downloading package averaged perceptron tagger eng to
     [nltk data]
                  /root/nltk_data...
     [nltk_data] Package averaged_perceptron_tagger_eng is already up-to-
     [nltk data]
                      date!
     [nltk_data] Downloading package conll2000 to /root/nltk_data...
     [nltk data] Package conll2000 is already up-to-date!
     [nltk_data] Downloading package movie_reviews to /root/nltk_data...
     [nltk_data] Package movie_reviews is already up-to-date!
     Finished.
# Import all necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.feature extraction.text import TfidfVectorizer
from sklearn.linear model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification report, confusion matrix, accuracy score
from sklearn.preprocessing import StandardScaler
from vaderSentiment.vaderSentiment import SentimentIntensityAnalyzer
from textblob import TextBlob
import librosa
import soundfile as sf
from IPython.display import Audio, display
import warnings
warnings.filterwarnings('ignore')
# Set up plotting style
plt.style.use('default')
sns.set_palette("husl")
print(" ✓ All libraries installed and imported successfully!")
print(" Ready to start building emotion detection systems!")
 → ✓ All libraries installed and imported successfully!
     Ready to start building emotion detection systems!
# Example text that demonstrates concepts from all previous modules
example_text = "I absolutely LOVE this amazing product! It works perfectly and exceeded my expectations."
print(" Analyzing with Previous Module Concepts:")
print(f"Original text: {example text}")
print()
# Module 2: Text Preprocessing
import re
import string
def basic_preprocess(text):
    # Tokenization (Module 2)
    tokens = text.split()
```

```
# Normalization (Module 2)
    tokens = [token.lower().strip(string.punctuation) for token in tokens]
    return tokens
tokens = basic preprocess(example text)
print(f" > Module 2 (Preprocessing) - Tokens: {tokens}")
# Module 4: Text Representation (TF-IDF preview)
sample docs = [
    "I love this product".
    "This product is amazing",
    "I hate this terrible product"
vectorizer = TfidfVectorizer()
tfidf matrix = vectorizer.fit transform(sample docs)
feature names = vectorizer.get feature names out()
print(f"\n !! Module 4 (Text Representation) - TF-IDF Features: {list(feature_names)}")
# Module 5: POS Tagging for sentiment
from textblob import TextBlob
blob = TextBlob(example_text)
pos tags = blob.tags
print(f"\n 			 Module 5 (POS Tagging) - Tags: {pos tags}")
print(" Notice: Adjectives (JJ) often carry emotional weight: 'amazing', 'perfect'")
# Module 6: Syntax for sentiment scope
print(f"\n Module 6 (Syntax/Parsing) - Dependency relationships help us understand:")
print(" - What is being loved? (the product)")
print(" - What makes it amazing? (it works perfectly)")
print(" - Scope of sentiment: entire product experience")
print("\n ♥ Key Insight: Sentiment analysis combines ALL previous concepts!")
print(" - Preprocessing cleans the text")
print(" - POS tags identify emotional words")
print(" - Syntax shows what the emotions target")
print(" - Text representation enables ML approaches")
→ Analyzing with Previous Module Concepts:
    Original text: I absolutely LOVE this amazing product! It works perfectly and exceeded my expectations.
     🍃 Module 2 (Preprocessing) - Tokens: ['i', 'absolutely', 'love', 'this', 'amazing', 'product', 'it', 'works', 'perfectly', 'and', 'exceeded', 'my', 'expectations']
     Module 4 (Text Representation) - TF-IDF Features: ['amazing', 'hate', 'is', 'love', 'product', 'terrible', 'this']
     Module 5 (POS Tagging) - Tags: [('I', 'PRP'), ('absolutely', 'RB'), ('LOVE', 'VBP'), ('this', 'DT'), ('amazing', 'JJ'), ('product', 'NN'), ('It', 'PRP'), ('works', 'VBZ'), ('perfectly', 'RB'),
       Notice: Adjectives (JJ) often carry emotional weight: 'amazing', 'perfect'
     ⚠ Module 6 (Syntax/Parsing) - Dependency relationships help us understand:
       - What is being loved? (the product)
        - What makes it amazing? (it works perfectly)
       - Scope of sentiment: entire product experience
     - Preprocessing cleans the text
       - POS tags identify emotional words
       - Syntax shows what the emotions target
        - Text representation enables ML approaches
```

Let's Recap some of the foundation from EDA & Machine Learning DataFrames Quick Review

What is a DataFrame?

Think of a DataFrame as a digital spreadsheet or table, just like Excel. It has:

- Rows (horizontal) each row represents one piece of data
- Columns (vertical) each column represents a different attribute or feature
- Cells where specific data values are stored

In Python, we use pandas (imported as pd) to create and work with DataFrames.

Step-by-Step Breakdown of Our Sample Data Creation:

Step 1: Creating the Text Data

```
sample_texts = [
    "I absolutely love this product! It's amazing and works perfectly.",
    "This is the worst purchase I've ever made. Completely disappointed.",
    # ... more examples
]
```

What's happening here:

- · We create a list (think of it as a container) of text examples
- · Each text represents what a real customer might write in a review or social media post
- The square brackets [] define a list in Python
- Each text is in quotes because it's a **string** (text data type)

Why these specific texts:

- Positive examples: "I absolutely love this product!" (clearly happy/positive)
- Negative examples: "This is the worst purchase!" (clearly angry/negative)
- Neutral examples: "The product is okay" (neither positive nor negative)
- Social media style: "OMG this is INCREDIBLE!!! 😅" (with emojis and excitement)

Step 2: Creating the DataFrame Structure

```
df_sample = pd.DataFrame({
    'text': sample_texts,
    'source': ['review', 'review', 'social_media', ...]
})
```

What's happening here:

- pd.DataFrame({...}) creates a new table/spreadsheet
- The curly braces {} define a dictionary a way to pair names with values
- 'text': sample_texts creates a column called "text" and fills it with our text examples
- 'source': [...] creates a column called "source" and labels each text as either "review" or "social_media"

The resulting table looks like this:

Index	text	source
0	"I absolutely love this product!"	review

Index	text	source
1	"This is the worst purchase"	review
2	"The product is okay"	review
3	"OMG this is INCREDIBLE!!!"	social_media
4	"Meh it's fine I guess"	review

6 Key Concepts Revisited:

- 1. Lists []: Store multiple items in order
- 2. **Dictionaries** { } : Pair names with values (like "name": "John")
- 3. DataFrames: Tables for organizing data with rows and columns
- 4. Columns: Vertical data categories (like "text" and "source")
- 5. Rows: Horizontal data entries (each customer comment)

Why Create This Sample Data?

1. Controlled Testing Environment

- We know exactly what emotions each text should express
- · We can test if our emotion detection systems work correctly
- · Like having an answer key to check our work

2. Diverse Text Types

- Product reviews: Professional, detailed feedback
- · Social media posts: Casual, with emojis and slang
- This tests how well our systems handle different writing styles

3. Balanced Emotions

- Mix of positive, negative, and neutral sentiments
- · Ensures our system can detect all types of emotions
- Prevents bias toward one emotion type

4. Real-World Simulation

- These texts represent what you'd actually find online
- · Helps us understand how emotion detection works in practice
- Prepares us for analyzing real customer feedback

Why This Matters for Emotion Analysis:

This sample data serves as our training ground where we can:

- · Test different emotion detection methods
- · Compare their accuracy
- Understand their strengths and weaknesses
- · Learn how to handle different types of text

Think of it like practicing basketball shots on a practice court before playing in a real game - we need controlled examples before analyzing real customer feedback!

Understanding the Output Code:

```
print(f"Total texts: {len(df_sample)}")
print(f"Sources: {df_sample['source'].value_counts().to_dict()}")
```

What each line does:

- len(df_sample) counts how many rows (text examples) we have
- df_sample['source'].value_counts() counts how many of each source type we have
- .to_dict() converts the count into a dictionary format for easy reading

This gives us a quick summary of our data structure and helps us verify everything was created correctly

```
# Create sample data that demonstrates different linguistic phenomena
sample texts = [
   "I absolutely love this product! It's amazing and works perfectly.", # Positive with intensifiers
    "This is the worst purchase I've ever made. Completely disappointed.", # Negative with superlatives
    "The product is okay, nothing special but does the job.", # Neutral with mixed signals
   "OMG this is INCREDIBLE!!!  Best thing ever!!!", # Social media style
   "I don't hate this product, but it's not great either.", # Negation (Module 6 syntax!)
   "The customer service was not bad, actually quite helpful.", # Double negation
   "This expensive product should work better for the price.", # Implicit criticism
    "Fast shipping, good quality, would recommend to others.", # Multiple aspects
   "It's fine I guess... could be worse "", # Lukewarm sentiment
    "TERRIBLE quality!!! Waste of money 😨 Never buying again!", # Strong negative
# Create DataFrame
df sample = pd.DataFrame({
    'text': sample_texts,
    'linguistic feature': [
        'intensifiers', 'superlatives', 'mixed_signals', 'social_media',
        'negation', 'double negation', 'implicit', 'multi aspect',
        'lukewarm', 'strong negative'
})
print(" > Sample Data with Linguistic Challenges:")
for i, row in df sample.iterrows():
   print(f"{i+1}. [{row['linguistic feature']}] {row['text']}") # Referring to the 'linguistic feature' column
print("\n > Sample Data Created:")
print(f"Total texts: {len(df_sample)}")
print("\n  First few examples:")
for i, row in df sample.head(3).iterrows():
   print(f"{i+1}. [{row['linguistic feature']}] {row['text']}")
 → Sample Data with Linguistic Challenges:
    1. [intensifiers] I absolutely love this product! It's amazing and works perfectly.
    2. [superlatives] This is the worst purchase I've ever made. Completely disappointed.
    3. [mixed signals] The product is okay, nothing special but does the job.
    4. [social_media] OMG this is INCREDIBLE!!! 👺 Best thing ever!!!
    5. [negation] I don't hate this product, but it's not great either.
    6. [double negation] The customer service was not bad, actually quite helpful.
    7. [implicit] This expensive product should work better for the price.
    8. [multi aspect] Fast shipping, good quality, would recommend to others.
    9. [lukewarm] It's fine I guess... could be worse 😜
    10. [strong negative] TERRIBLE quality!!! Waste of money 🔯 Never buying again!
     Sample Data Created:
```

```
Total texts: 10

A First few examples:

1. [intensifiers] I absolutely love this product! It's amazing and works perfectly.

2. [superlatives] This is the worst purchase I've ever made. Completely disappointed.

3. [mixed_signals] The product is okay, nothing special but does the job.
```

Part 1: Text Sentiment Analysis with VADER and TextBlob

Understanding Rule-Based Sentiment Analysis

Rule-based sentiment analysis uses **lexicons** (dictionaries) that map words to sentiment scores. This connects directly to Module 4's discussion of text representation - instead of TF-IDF vectors, we use pre-built emotional dictionaries.

VADER (Valence Aware Dictionary and sEntiment Reasoner):

- · Designed for social media text
- Handles emojis, capitalization, punctuation
- Uses linguistic rules (like Module 6's parsing rules)

TextBlob:

- General-purpose sentiment analysis
- · Provides both polarity and subjectivity scores
- · Based on movie review corpus

→ 1.1 VADER Sentiment Analysis

VADER (Valence Aware Dictionary and sEntiment Reasoner) is specifically designed for social media text. It:

- · Handles emojis, slang, and intensifiers
- Provides compound scores from -1 (most negative) to +1 (most positive)
- Gives detailed breakdowns (positive, negative, neutral)

Let's see how VADER analyzes our sample texts:

```
# Initialize VADER analyzer
vader_analyzer = SentimentIntensityAnalyzer()
def analyze_with_vader(text):
   """Analyze sentiment using VADER"""
   scores = vader_analyzer.polarity_scores(text)
   return scores
def classify_vader_sentiment(compound_score):
   """Convert VADER compound score to sentiment label"""
   if compound score >= 0.05:
       return 'Positive'
   elif compound score <= -0.05:
       return 'Negative'
   else:
       return 'Neutral'
# Analyze all sample texts with VADER
vader_results = []
for text in df_sample['text']:
   scores = analyze_with_vader(text)
```

```
vader_results.append({
        'compound': scores['compound'],
        'positive': scores['pos'],
        'negative': scores['neg'],
        'neutral': scores['neu'],
        'sentiment': classify vader sentiment(scores['compound'])
   })
# Add VADER results to DataFrame
vader df = pd.DataFrame(vader results)
df_sample = pd.concat([df_sample, vader_df.add_prefix('vader_')], axis=1)
print("@ VADER Analysis Results:")
print(df_sample['vader_sentiment'].value_counts())
print("\n Q Detailed Results:")
for i, row in df_sample.iterrows():
   print(f"{i+1}. {row['vader_sentiment']} (score: {row['vader_compound']:.3f})")
    print(f" Text: {row['text'][:60]}...")
    print()
→ * VADER Analysis Results:
     Sentiment Distribution:
     vader sentiment
     Positive 5
     Negative 4
     Neutral 1
    Name: count, dtype: int64
     Q Detailed Results:
    1. Positive (score: 0.930)
        Text: I absolutely love this product! It's amazing and works perfe...
    2. Negative (score: -0.817)
        Text: This is the worst purchase I've ever made. Completely disapp...
    3. Neutral (score: -0.046)
        Text: The product is okay, nothing special but does the job....
    4. Positive (score: 0.886)
       Text: OMG this is INCREDIBLE!!! 👺 Best thing ever!!!...
    5. Negative (score: -0.533)
        Text: I don't hate this product, but it's not great either....
    6. Positive (score: 0.713)
        Text: The customer service was not bad, actually quite helpful....
    7. Positive (score: 0.440)
        Text: This expensive product should work better for the price....
    8. Positive (score: 0.660)
       Text: Fast shipping, good quality, would recommend to others....
    9. Negative (score: -0.318)
        Text: It's fine I guess... could be worse ⊕...
     10. Negative (score: -0.832)
        Text: TERRIBLE quality!!! Waste of money ☑ Never buying again!...
```

TextBlob provides a different approach to sentiment analysis:

- · Uses a different lexicon and algorithm than VADER
- Provides polarity (-1 to +1) and subjectivity (0 to 1) scores
- Generally more conservative in its sentiment assignments

Let's compare TextBlob results with VADER:

```
# TextBlob Analysis
def analyze_with_textblob(text):
   """Analyze sentiment using TextBlob"""
   blob = TextBlob(text)
   return {
       'polarity': blob.sentiment.polarity,
       'subjectivity': blob.sentiment.subjectivity
def classify_textblob_sentiment(polarity):
   """Convert TextBlob polarity to sentiment label"""
   if polarity > 0.1:
       return 'Positive'
   elif polarity < -0.1:
       return 'Negative'
   else:
       return 'Neutral'
# Analyze with TextBlob
textblob results = []
for text in df_sample['text']:
   scores = analyze_with_textblob(text)
   textblob results.append({
       'polarity': scores['polarity'],
       'subjectivity': scores['subjectivity'],
       'sentiment': classify_textblob_sentiment(scores['polarity'])
   })
# Add TextBlob results
textblob_df = pd.DataFrame(textblob_results)
df sample = pd.concat([df sample, textblob df.add prefix('textblob ')], axis=1)
print("@ VADER vs TextBlob Comparison:")
comparison = pd.crosstab(df_sample['vader_sentiment'], df_sample['textblob_sentiment'], margins=True)
print(comparison)
print("\n Q Detailed Comparison:")
for i, row in df_sample.iterrows():
   if row['vader_sentiment'] != row['textblob_sentiment']:
       print(f"DISAGREEMENT on [{row['linguistic feature']}]:")
       print(f" Text: {row['text']}")
       print(f" VADER: {row['vader_sentiment']} ({row['vader_compound']:.3f})")
       print(f" TextBlob: {row['textblob_sentiment']} ({row['textblob_polarity']:.3f})")
       print()
print("@ TextBlob Analysis Results:")
print(df_sample['textblob_sentiment'].value_counts())
comparison = pd.crosstab(df_sample['vader_sentiment'], df_sample['textblob_sentiment'], margins=True)
print(comparison)
```

```
textblob sentiment Negative Neutral Positive All
   vader_sentiment
   Negative
                                        0 4
                                1
   Neutral
                         0
                                0
                                        1 1
   Positive
                         0
                                1
                                        4 5
   All
   Detailed Comparison:
   DISAGREEMENT on [mixed signals]:
    Text: The product is okay, nothing special but does the job.
    VADER: Neutral (-0.046)
    TextBlob: Positive (0.429)
   DISAGREEMENT on [implicit]:
    Text: This expensive product should work better for the price.
    VADER: Positive (0.440)
    TextBlob: Neutral (0.000)
   DISAGREEMENT on [lukewarm]:
    Text: It's fine I guess... could be worse ≅
    VADER: Negative (-0.318)
    TextBlob: Neutral (0.008)
   Sentiment Distribution:
   textblob sentiment
   Positive 5
   Negative 3
   Neutral
   Name: count, dtype: int64
   Comparison with VADER:
   textblob_sentiment Negative Neutral Positive All
   vader_sentiment
   Negative
                        3
                                1
                                        0 4
                     0 0 1
                            0
                                        1 1
   Neutral
   Positive
                                        4 5
```

✓ YOUR TURN - Exercise 1: Analyzing Different Text Types

Now it's your turn to experiment with VADER and TextBlob! Complete the following tasks:

```
# / YOUR TURN: Create examples from YOUR experience

# 1. Think of 5 real examples from YOUR life where sentiment might be ambiguous
# These should be based on:
# - Something you actually said or wrote
# - A review you read that confused you
# - Social media posts from your friends/family
# - Text messages where tone was misunderstood
# - Cultural expressions from your background

your_real_examples = [
# TODO: Replace with YOUR actual examples
    "Example 1: How sweet are the customers at the fiesta bakery? ",
    "Example 2: Chick fil a: This chick fil a has a warm and inviting environment, but the food was cold. ", #started off positive, ended negative
    "Example 3: a post my cousin made was (You stole my post, Im coming for your catalytic converter next my guy) ", #sarcasm
    "Example 4: Bad news, my car is going to the shop for new tint tomorrow, so I will not be able to make it.",
    "Example 5: No manches!, Bless your heart! All talk, no bite!"

]
```

```
# 2. Analyze YOUR examples
print("  Analysis of YOUR Real-World Examples:")
print("=" * 50)
for i, text in enumerate(your real examples, 1):
    # TODO: Analyze each with VADER and TextBlob
    vader_scores = analyze_with_vader(text)
    textblob scores = analyze with textblob(text)
    print(f"\n{i}. Your Example: {text}")
    print(f" VADER: {classify_vader_sentiment(vader_scores['compound'])} ({vader_scores['compound']:.3f})")
    print(f" TextBlob: {classify textblob sentiment(textblob scores['polarity'])} ({textblob scores['polarity']:.3f})")
# 3. Reflection Questions (AI cannot answer these for you!)
print("\n@ Personal Reflection Questions:")
print("Answer these based on YOUR experience running the code above:")
print()
print("1. Which of YOUR examples did the algorithms get wrong? Why do you think that happened?")
print(" : The algorithms did not give the right results for example 4. I think this occured because it associated 'bad news' with a negative tone. Rather, it was good because my boss is getting new 1
print("2. Did any results surprise you? Which ones and why?")
print(" Your answer: The result for example 5 surprised me. It gave the right result and I used spanish/english in the same line. This was surprising!")
print()
print("3. How did your cultural background or personal communication style affect the results?")
print(" Your answer: My cultural background example did not seem to affect the results.")
print()
print("4. If you were to improve these algorithms, what would you focus on based on YOUR examples?")
print(" Your answer: I would add strong positives since we added strong negative. This would probably help with my results for my examples. ")
     Analysis of YOUR Real-World Examples:
     1. Your Example: Example 1: How sweet are the customers at the fiesta bakery?
        VADER: Positive (0.727)
        TextBlob: Positive (0.350)
     2. Your Example: Example 2: Chick fil a: This chick fil a has a warm and inviting environment, but the food was cold.
        VADER: Positive (0.273)
        TextBlob: Neutral (0.000)
     3. Your Example: Example 3: a post my cousin made was (You stole my post, Im coming for your catalytic converter next my guy)
        VADER: Neutral (0.000)
        TextBlob: Neutral (0.000)
     4. Your Example: Example 4: Bad news, my car is going to the shop for new tint tomorrow, so I will not be able to make it.
        VADER: Negative (-0.542)
        TextBlob: Neutral (-0.021)
5. Your Example: Example 5: No manches!, Bless your heart! All talk, no bite!
        VADER: Negative (-0.342)
        TextBlob: Neutral (0.000)
     Personal Reflection Ouestions:
     Answer these based on YOUR experience running the code above:
     1. Which of YOUR examples did the algorithms get wrong? Why do you think that happened?
        : The algorithms did not give the right results for example 4. I think this occured because it associated 'bad news' with a negative tone. Rather, it was good because my boss is getting new tin
     2. Did any results surprise you? Which ones and why?
        Your answer: The result for example 5 surprised me. It gave the right result and I used spanish/english in the same line. This was surprising!
     3. How did your cultural background or personal communication style affect the results?
        Your answer: My cultural background example did not seem to affect the results.
```

4. If you were to improve these algorithms, what would you focus on based on YOUR examples?

Your answer: I would add strong positives since we added strong negative. This would probably help with my results for my examples.

🗸 🔎 Conceptual Understanding Check

Answer these questions to cement your learning:

Q1: How does sentiment analysis connect to Module 6 (Syntax/Parsing)?

Your answer: Sentiment analysis and syntax/parsing are connected by their focus on relationships between words in a sentence/paragraph.

Q2: Why might POS tags from Module 5 be useful for sentiment analysis?

Your answer: POS tags break down words into the different parts of speech which can help sentiment analysis by not putting the same word into the same context. It enables the ability to differentiate between word meaning and the message needing to be recieved in a positive tone rather than negative.

Q3: How does text preprocessing from Module 2 affect sentiment analysis accuracy?

Your answer: Text preprocessing is a huge part of the backbone for sentiment analysis. After it cleans and converts the raw text, it becomes ready for the analysis stage!

Q4: Compare rule-based sentiment analysis to the TF-IDF approach from Module 4. What are the key differences?

Your answer: Key differences: 1.Rule-based sentiment requires far more manual effort to maintain the rules. TF-IDF is based more on frequency of terms in the deliverable.

Part 2: Machine Learning Approach to Sentiment Analysis

Moving Beyond Rule-Based Methods

While rule-based methods like VADER and TextBlob are excellent for general-purpose sentiment analysis, a machine learning approach offers several key advantages. By training a model on a specific dataset, we can:

- Learn domain-specific patterns from your data
- Handle context better through feature learning
- Adapt to specific use cases like product reviews vs. social media
- · Potentially achieve higher accuracy on your specific dataset

To use a machine learning algorithm for a task like sentiment classification, we must first convert our text data into a numerical format that the model can understand. This process is called feature engineering or text representation.

Let's build a machine learning classifier using scikit-learn and compare it with our rule-based methods.

Loading Data for Machine Learning

In a real-world machine learning sentiment analysis project, you would typically load a much larger dataset with a minimum of thousands or even millions of tokens. This dataset would likely be loaded from a file (like a CSV or JSON) or an API. The larger the dataset, the better the model can learn complex patterns and generalize to new, unseen text.

Note: For the purpose of this exercise, we will be using a short, illustrative text dataset to demonstrate the process.

```
# Create ML dataset with examples that test different linguistic phenomena ml texts = \lceil
```

```
# Positive examples with different structures
    "Being awarded the neighborhood exceeded my expectations! I am proud of the team.", # Simple positive
    "The new brand of tortillas are amazing! They taste homemade, we will purchase again!", # Multiple positive aspects
    "You are amazing and I cannot wait to see you! I am ready for you to jumpp into my arms!", # Enthusiastic
    "Not bad at all, it was a great deal for what we got.", # Positive via negation
    "Honestly, I did not expect it to be as good as the reviews said. It was truly some of the best food around.", # Exceeded expectations
    "Outstanding quality of video. Very satisfied.", # Value-based positive
    "Going into this, I knew what I wanted. Thankfully I received the best outcome!", # Expectation match
    "Incredible skills and performace.", # Superlatives
    "Couldn't be happier with anyone else in the world.", # Negative construction, positive meaning
    "You did great! Thank you for your hardwork.", # Multiple positive attributes
    # Negative examples with different structures
    "Terrible quality. Made with the cheapest material.", # Simple negative
    "Worst customer service ever. Very unhappy with this location.", # Superlative negative
    "Complete waste of money. Will never travel to come here again. Waste of time.", # Strong negative advice
    "Not what I expected, made me quite angry when I opened the package", # Expectation mismatch
    "Arrived damaged and customer service would not refund me.", # Multiple negative aspects
    "Doesn't work as advertised. Very wobbly and unreliable.", # Functional failure
    "Cheap material that will not perform its function properly.", # Quality criticism
    "Overpriced for the amount of food you get", # Value-based negative
    "Regret buying this. Made me sad to know I cannot get my money back.", # Emotional regret
    "Horrible experience with the ride restrictions. Not catered to short people.", # Comprehensive negative
    # Neutral examples with different structures
    "The product is okay. Nothing to make it shine.", # Simple neutral
    "Average quality for what you expect from a fast food place.", # Expectation match
    "It works but I would not be able to use it in tight places". # Functional but limited
    "Decent product, meets the standard of the market.", # Adequate performance
    "Nothing special, just average.", # Explicit neutrality
    "Standard quality, we were not expecting much more.", # Expectation match
    "It's okay for what we need it for today.", # Conditional acceptance
    "Acceptable quality, nothing special", # Adequate but unremarkable
    "Does the job but nothing to rave about.", # Functional adequacy
    "Mediocre product with simple outcome." # Below average but functional
ml labels = (
    ['positive'] * 10 +
    ['negative'] * 10 +
    ['neutral'] * 10
# Create DataFrame
df ml = pd.DataFrame({
    'text': ml texts,
    'sentiment': ml labels
print(f"Total examples: {len(df ml)}")
print(f"Label distribution: {df ml['sentiment'].value counts().to dict()}")
# Show connection to Module 4 concepts
print("- Each text will become a TF-IDF vector (like Module 4)")
print("- But now we have LABELS (sentiment) for supervised learning")
print("- This transforms unsupervised representation into supervised classification")
 → Machine Learning Dataset:
     Total examples: 30
     Label distribution: {'positive': 10, 'negative': 10, 'neutral': 10}
```

```
    ℰ Connection to Module 4 (Text Representation):
    Each text will become a TF-IDF vector (like Module 4)
    But now we have LABELS (sentiment) for supervised learning
    This transforms unsupervised representation into supervised classification
```

2.1 Feature Extraction with TF-IDF

Remember before we can train a machine learning model, we need to convert our text into numerical features. We'll use TF-IDF (Term

Frequency-Inverse Document Frequency) which:

- · Captures the importance of words in documents
- Reduces the impact of common words
- Creates sparse but meaningful feature vectors

```
# Feature extraction using TF-IDF (Module 4 concepts)
# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(
    df ml['text'],
    df ml['sentiment'],
    test size=0.3,
    random state=42,
    stratify=df_ml['sentiment']
# Create TF-IDF vectorizer
tfidf vectorizer = TfidfVectorizer(
   max_features=1000,  # Limit vocabulary size
    stop_words='english', # Remove common English stop words
    ngram_range=(1, 2), # Unigrams and bigrams
    lowercase=True
# Fit the vectorizer on training data and transform both sets
X_train_tfidf = tfidf_vectorizer.fit_transform(X_train)
X test tfidf = tfidf vectorizer.transform(X test)
print("    TF-IDF Feature Extraction (Module 4 Review):")
print(f"Training matrix shape: {X_train_tfidf.shape}")
print(f"Vocabulary size: {len(tfidf_vectorizer.vocabulary_)}")
print(f"Feature sparsity: \{(1 - X_train_tfidf.nnz / (X_train_tfidf.shape[\emptyset] * X_train_tfidf.shape[1])):.3f\}")
# Show some features
feature_names = tfidf_vectorizer.get_feature_names_out()
print(f"\nSample features: {list(feature_names[:15])}")
 TF-IDF Feature Extraction (Module 4 Review):
     Training matrix shape: (21, 157)
     Vocabulary size: 157
     Feature sparsity: 0.949
     Sample features: ['acceptable', 'acceptable quality', 'advertised', 'advertised wobbly', 'amazing', 'amazing wait', 'arms', 'average', 'average quality', 'awarded', 'awarded neighborhood', 'bad',
```

2.2 Training Multiple Classifiers

Let's train and compare different machine learning algorithms:

• Logistic Regression: Linear model, fast and interpretable

• Random Forest: Ensemble method, handles non-linear patterns

We'll evaluate their performance and see how they compare to our rule-based methods.

```
# Train Logistic Regression
lr model = LogisticRegression(random state=42, max iter=1000)
lr_model.fit(X_train_tfidf, y_train)
# Train Random Forest
rf model = RandomForestClassifier(n estimators=100, random state=42)
rf_model.fit(X_train_tfidf, y_train)
# Make predictions
lr_predictions = lr_model.predict(X_test_tfidf)
rf_predictions = rf_model.predict(X_test_tfidf)
# Calculate accuracies
lr_accuracy = accuracy_score(y_test, lr_predictions)
rf_accuracy = accuracy_score(y_test, rf_predictions)
print("@ Model Performance:")
print(f"Logistic Regression Accuracy: {lr accuracy:.3f}")
print(f"Random Forest Accuracy: {rf accuracy:.3f}")
# Detailed classification reports
print(classification_report(y_test, lr_predictions))
print("\n | Random Forest Classification Report:")
print(classification_report(y_test, rf_predictions))
Logistic Regression Accuracy: 0.778
    Random Forest Accuracy: 0.556
    Logistic Regression Classification Report:
                precision recall f1-score support
       negative
                    0.60
                          1.00
                                     0.75
                                                 3
        neutral
                   1.00 0.67
                                     0.80
                                                 3
       positive
                  1.00 0.67
                                     0.80
                                                 3
                                     0.78
       accuracy
       macro avg
                    0.87 0.78
                                     0.78
                                                 9
    weighted avg
                    0.87 0.78
                                     0.78
                                                 9
    Random Forest Classification Report:
                precision recall f1-score support
       negative
                    0.43
                          1.00
                                     0.60
                                                 3
        neutral
                    1.00 0.67
                                     0.80
                                                 3
       positive
                    0.00 0.00
                                     0.00
                                                 3
                                                 9
                                     0.56
       accuracy
       macro avg
                    0.48 0.56
                                     0.47
                                                 9
    weighted avg
                    0.48 0.56
                                     0.47
                                                 9
```

2.3 Model Interpretation

Let's understand what our models have learned by examining the most important features for each sentiment class.

```
# Get feature importance from Logistic Regression
def get_top_features(model, vectorizer, class_labels, n_features=5):
    """Extract top features for each class from a trained model"""
   feature names = vectorizer.get feature names out()
   for i, class label in enumerate(class labels):
       if hasattr(model, 'coef '):
           # For linear models like Logistic Regression
           coefficients = model.coef [i]
            top indices = coefficients.argsort()[-n features:][::-1]
            top_features = [(feature_names[idx], coefficients[idx]) for idx in top_indices]
       else:
            # For tree-based models like Random Forest
            importances = model.feature_importances_
            top indices = importances.argsort()[-n features:][::-1]
            top features = [(feature names[idx], importances[idx]) for idx in top indices]
        print(f"\nQ Top features for '{class_label}':")
        for feature, score in top features:
           print(f" {feature}: {score:.3f}")
# Analyze Logistic Regression features
print("   Logistic Regression - Most Important Features:")
get_top_features(lr_model, tfidf_vectorizer, lr_model.classes_)
# Analyze Random Forest features (overall importance)
print("\n \( \bar{\Pi} \) Random Forest - Most Important Features (Overall):")
feature names = tfidf vectorizer.get feature names out()
importances = rf_model.feature_importances_
top indices = importances.argsort()[-10:][::-1]
print("Top 10 most important features:")
for idx in top indices:
   print(f" {feature_names[idx]}: {importances[idx]:.3f}")
→ Logistic Regression - Most Important Features:
     Top features for 'negative':
      overpriced food: 0.297
      overpriced: 0.297
      monev: 0.250
      waste: 0.247
      terrible: 0.198
     Top features for 'neutral':
      special: 0.391
      standard: 0.344
      average: 0.326
      product: 0.306
      quality: 0.276
     Top features for 'positive':
      great: 0.318
      happier world: 0.217
      happier: 0.217
      couldn: 0.217
      world: 0.217
     Random Forest - Most Important Features (Overall):
     Top 10 most important features:
      quality: 0.035
      special: 0.027
      average: 0.023
      standard: 0.021
```

```
money: 0.019
overpriced: 0.017
happier: 0.017
great: 0.017
does job: 0.017
overpriced food: 0.016
```

✓ YOUR TURN - Exercise 2: Comparing All Methods

Now let's compare all our sentiment analysis methods on the same test data!

```
# ℰ YOUR TURN: Compare all methods on the same data
# 1. Apply the same test set to all models
test_texts = X_test.tolist()
true_labels = y_test.tolist()
# TODO: Get predictions from all methods
# (Assumes lr predictions and rf predictions are already available from the previous step)
vader predictions = []
textblob_predictions = []
# TODO: For each text in test_texts, get predictions from VADER and TextBlob
for text in test texts:
   # VADER prediction
   vader_score = analyze_with_vader(text)['compound']
   vader_pred = classify_vader_sentiment(vader_score).lower()
   vader_predictions.append(vader_pred)
   # TextBlob prediction
   textblob score = analyze with textblob(text)['polarity']
   textblob_pred = classify_textblob_sentiment(textblob_score).lower()
   textblob_predictions.append(textblob_pred)
# 2. Calculate final accuracies
vader acc = accuracy score(true labels, vader predictions)
textblob_acc = accuracy_score(true_labels, textblob_predictions)
lr_acc = accuracy_score(true_labels, lr_predictions)
rf_acc = accuracy_score(true_labels, rf_predictions)
# 3. Print the high-level comparison results
print("=" * 40)
                                {vader_acc:.3f}")
print(f"VADER Accuracy:
print(f"TextBlob Accuracy:
                                {textblob acc:.3f}")
print(f"Logistic Regression:
                               {lr acc:.3f}")
print(f"Random Forest:
                               {rf acc:.3f}")
# 4. Analyze and print specific disagreements
print("\n \ Analyzing Disagreements in the Results:")
print("=" * 40)
disagreements = 0
for i, (text, true_label) in enumerate(zip(test_texts, true_labels)):
   predictions = {
        'VADER': vader_predictions[i],
       'TextBlob': textblob_predictions[i],
       'LogReg': lr predictions[i],
        'RandomForest': rf predictions[i]
```

```
# Check if there was any disagreement among the predictions
   if len(set(predictions.values())) > 1:
       disagreements += 1
       print(f"\nDisagreement #{disagreements}:")
       print(f"Text:
                          {text}")
       print(f"True Label: {true label}")
       for method, pred in predictions.items():
           correct = "\" if pred == true_label else "X"
           print(f" -> {method:<15} predicted: {pred:<10} {correct}")</pre>
print(f"\n\n Total disagreements found among the models: {disagreements}")
Disagreement #2:
    Text:
            Incredible skills and performace.
    True Label: positive
      -> VADER
                        predicted: neutral X
      -> TextBlob
                        predicted: positive ✓
                        predicted: negative X
      -> LogReg
      -> RandomForest predicted: negative X
    Disagreement #3:
    Text: The new brand of tortillas are amazing! They taste homemade, we will purchase again!
    True Label: positive
      -> VADER
                        predicted: positive ✓
      -> TextBlob
                        predicted: positive ✓
      -> LogReg
                        predicted: positive
      -> RandomForest predicted: negative X
    Disagreement #4:
    Text: It's okay for what we need it for today.
    True Label: neutral
      -> VADER
                        predicted: positive X
                        predicted: positive
                                            X
      -> TextBlob
      -> LogReg
                        predicted: neutral
      -> RandomForest predicted: neutral
    Disagreement #5:
            It works but I would not be able to use it in tight places
    Text:
    True Label: neutral
      -> VADER
                        predicted: neutral ✓
      -> TextBlob
                        predicted: positive X
      -> LogReg
                        predicted: negative
      -> RandomForest predicted: negative X
    Disagreement #6:
             Honestly, I did not expect it to be as good as the reviews said. It was truly some of the best food around.
    Text:
    True Label: positive
      -> VADER
                        predicted: positive ✓
      -> TextBlob
                        predicted: positive ✓
```

Hands-On Analysis Questions

Answer these based on YOUR specific results above:

Q1: Looking at YOUR results above, which method performed best? Was this what you expected?

Your answer based on your results:

Logistic regression performed the best. I did think that VADER would perform better though.

Q2: Examine the disagreements YOUR code found. Pick one disagreement and explain why you think the methods disagreed.

Your analysis of a specific disagreement:

Disagreement #8 showed that VADER, LogReg, and RandomForest all agreed, but TextBlob did not. I think it is due to it being semi informal. TextBlob performs well for more structured text.

Q3: How do the TF-IDF features (from Module 4) help the ML models in a way that the rule-based methods can't access?

Your answer explaining the role of TF-IDF: TF-IDF help ML models differently by being able to not follow strict manual rules. TF-IDF can look at the whole corpus to capture significance of words, even those not strictly defined.

Q4: If you had to choose ONE method for a real business application, which would you choose based on YOUR results? Why?

Your final decision and reasoning. A complete answer should consider the trade-offs by describing a scenario where a rule-based model like VADER might be better, and another scenario where the ML model is the only choice.

I would choose the LogReg for a business application. I would use it to help navigate spam emails due to the large volume of emails sent it to the company. Many people have access to the email addresses of our employees, and the statistical aspect of LogRef is what can allow this to happen. Unlike TextBlob which would not be useful for this task. Is is limited for small scale-training projects. It is also not great for domain-specific tasks.

Part 3: Speech Emotion Detection

From Audio Preprocessing to Emotion Analysish

In Module 3, you learned how to process and clean audio signals. Now, we'll take that a step further. Instead of just cleaning audio, we will extract emotional features from speech signals. The goal is to identify measurable characteristics of a sound wave that correspond to how an emotion is expressed vocally.

- What is being said? (Text Analysis Parts 1 & 2)
- How is it being said? (Speech Analysis Part 3)List item

We will use the librosa library to analyze audio and extract features like pitch, energy, and tempo, which are crucial for detecting emotion in speech.

Key Connection: Audio Preprocessing + Feature Extraction + Machine Learning = Speech Emotion Detection

Speech carries emotional information that text alone cannot capture:

- · Prosodic features: Pitch, rhythm, stress patterns
- · Voice quality: Tone, breathiness, tension
- Temporal dynamics: Speaking rate, pauses

Let's build a speech emotion detector using audio signal processing techniques.

Note: In this lab, we'll simulate audio features since we're working in a text-based environment. In a real application, you would extract these features from actual audio files

```
# In a real application, these would be extracted from audio files using librosa
# Create sample audio signals that demonstrate emotional speech patterns
def create_emotional_audio(emotion, duration=2, sample_rate=22050):
    """Create audio that simulates emotional speech patterns"""
   t = np.linspace(0, duration, int(sample_rate * duration))
    if emotion == 'happy':
       # Happy: Higher pitch, more variation, upward intonation
       base\_freq = 225
       pitch variation = 60 * np.sin(2 * np.pi * 2 * t)
       amplitude = 0.9
    elif emotion == 'sad':
       # Sad: Lower pitch, less variation, downward intonation
       base freq = 145
       pitch_variation = 11 * np.sin(2 * np.pi * 0.5 * t)
       amplitude = 0.4
    elif emotion == 'angry':
       # Angry: Variable pitch, harsh quality, higher energy
       base freq = 210
       pitch_variation = 90 * np.sin(2 * np.pi * 3 * t) * np.random.normal(1, 0.2, len(t))
       amplitude = 0.9
    else: # neutral
       # Neutral: Steady pitch, moderate energy
       base freq = 185
       pitch variation = 26 * np.sin(2 * np.pi * 1 * t)
       amplitude = 0.6
    # Create the audio signal
    frequency = base freq + pitch variation
    audio = amplitude * np.sin(2 * np.pi * frequency * t)
    # Add harmonics for more realistic sound
    audio += 0.4 * amplitude * np.sin(2 * np.pi * frequency * 2 * t)
    audio += 0.2 * amplitude * np.sin(2 * np.pi * frequency * 3 * t)
    # Add noise for realism
   noise = 0.05 * np.random.normal(0, 1, len(audio))
    audio += noise
   return audio, sample rate
# Create and play sample audio for each emotion
emotions = ['happy', 'sad', 'angry', 'neutral']
audio_samples = {}
print("Listen to how different emotions sound in speech patterns...\n")
for emotion in emotions:
    audio, sr = create emotional audio(emotion)
    audio samples[emotion] = (audio, sr)
    print(f" \{ \{ emotion.capitalize() \} Audio:")
    display(Audio(audio, rate=sr))
    print(f" Characteristics: {emotion} speech patterns simulated")
    print()
```

```
print(" ? Connection to Module 3:")
print("- We're using the same audio processing concepts (sample rate, frequency).")
print("- But now we are extracting EMOTIONAL features instead of just cleaning the audio.")
 Creating Emotional Speech Samples:
    Listen to how different emotions sound in speech patterns...
     🥞 Happy Audio:
           0:00 / 0:02
        Characteristics: happy speech patterns simulated
     😼 Sad Audio:
           0:00 / 0:02
        Characteristics: sad speech patterns simulated
     Mangry Audio:
           0:02 / 0:02
        Characteristics: angry speech patterns simulated
     Neutral Audio:
           0:00 / 0:02
        Characteristics: neutral speech patterns simulated
     P Connection to Module 3:
     - We're using the same audio processing concepts (sample rate, frequency).
     - But now we are extracting EMOTIONAL features instead of just cleaning the audio.
```

3.1 Extracting Emotional Features from Audio

Now that we have audio signals, we'll use librosa to extract features that quantify the emotional characteristics we just heard. These features include:

Pitch: The fundamental frequency of the voice (how high or low it is).

- Energy (RMS): The loudness or intensity of the speech.
- Tempo: The speed or pace of speech (Beats Per Minute).
- Spectral Centroid: Relates to the "brightness" of a sound.

```
# Extract emotional features from audio (building on Module 3)
def extract_emotional_features(audio, sample_rate):
    """Extract features that indicate emotional state"""
    features = {}

# Prosodic features (how we speak)
    pitches, magnitudes = librosa.piptrack(y=audio, sr=sample_rate)
    pitch_values = []
    for t in range(pitches.shape[1]):
        index = magnitudes[:, t].argmax()
        pitch = pitches[index, t]
```

```
if pitch > 0:
           pitch_values.append(pitch)
   if pitch values:
       features['pitch mean'] = np.mean(pitch values)
       features['pitch std'] = np.std(pitch values)
   else:
       features['pitch_mean'] = 0
       features['pitch_std'] = 0
   # Energy features
   rms_energy = librosa.feature.rms(y=audio)[0]
   features['energy_mean'] = np.mean(rms_energy)
   features['energy_std'] = np.std(rms_energy)
   # Spectral features
   spectral_centroids = librosa.feature.spectral_centroid(y=audio, sr=sample_rate)[0]
   features['spectral_centroid_mean'] = np.mean(spectral_centroids)
   # Tempo
   tempo, = librosa.beat.beat track(y=audio, sr=sample rate)
   features['tempo'] = tempo
   return features
# Extract features from our emotional audio samples
print("    Extracting Emotional Features from Audio:")
audio feature data = []
# Create multiple samples per emotion for better analysis
for emotion in emotions:
   print(f"\nProcessing {emotion.capitalize()} emotion samples...")
   for i in range(25): # 25 samples per emotion for a balanced dataset
       audio, sr = create_emotional_audio(emotion, duration=1.5)
       features = extract_emotional_features(audio, sr)
       features['emotion'] = emotion
       audio feature data.append(features)
# Create DataFrame for analysis
df audio = pd.DataFrame(audio feature data)
print(f"\n Audio emotion dataset created: {len(df_audio)} samples, {len(df_audio.columns)-1} features")
print("\nFirst 5 rows of the feature dataset:")
print(df audio.head())
Processing Happy emotion samples...
    Processing Sad emotion samples...
    Processing Angry emotion samples...
    Processing Neutral emotion samples...
     Audio emotion dataset created: 100 samples, 6 features
    First 5 rows of the feature dataset:
       pitch_mean pitch_std energy_mean energy_std spectral_centroid_mean \
    0 464.148615 285.785109 0.697412 0.042495
                                                                1875.711761
    1 464.031767 285.757665 0.697751 0.042765
                                                                1870.627813
    2 462.802145 283.880915 0.697379 0.042552
                                                                1877.732705
    3 458.668927 284.831706
                              0.697979 0.042644
                                                                1872.156243
    4 462.527061 284.619037
                                                                1876.782464
                              0.696994 0.042584
```

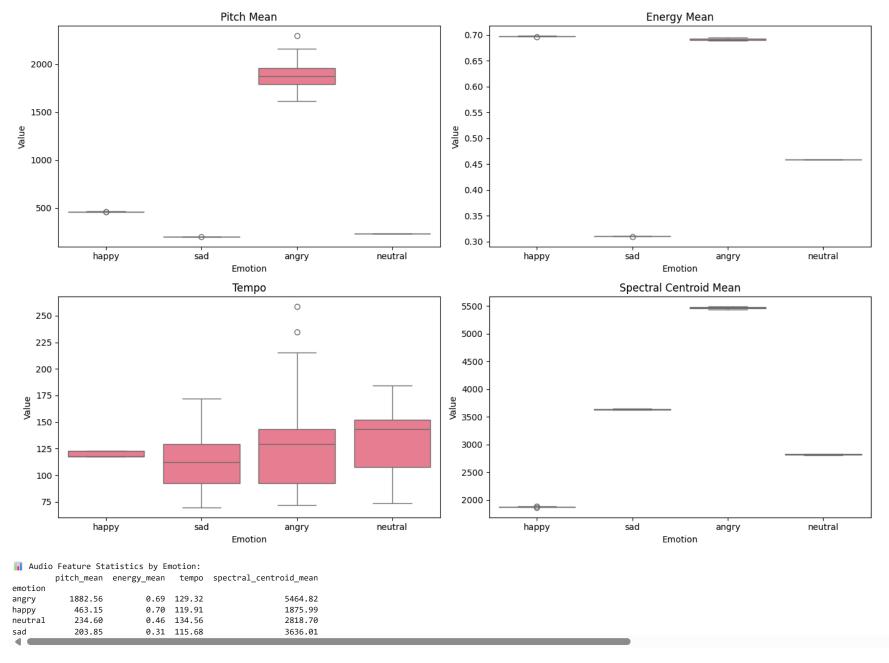
```
tempo emotion
0 [123.046875] happy
1 [123.046875] happy
2 [117.45383522727273] happy
3 [123.046875] happy
4 [117.45383522727273] happy
```

→ 3.2 Visualizing Audio Features

Before building a classifier, let's visualize the features we just extracted. This helps us see if there are clear, measurable differences between the emotions

```
# Visualize key audio features by emotion
fig, axes = plt.subplots(2, 2, figsize=(14, 10))
fig.suptitle('Audio Features by Emotion', fontsize=16)
key_features = ['pitch_mean', 'energy_mean', 'tempo', 'spectral_centroid_mean']
# Flatten the 'tempo' column if it contains arrays
df_audio['tempo'] = df_audio['tempo'].apply(lambda x: x[0] if isinstance(x, np.ndarray) else x)
for i, feature in enumerate(key_features):
   row, col = i // 2, i % 2
   sns.boxplot(x='emotion', y=feature, data=df_audio, ax=axes[row, col])
   axes[row, col].set_title(f'{feature.replace("_", " ").title()}')
   axes[row, col].set_xlabel('Emotion')
   axes[row, col].set_ylabel('Value')
plt.tight_layout(rect=[0, 0.03, 1, 0.95])
plt.show()
# Statistical summary
print("\n  Audio Feature Statistics by Emotion:")
summary_stats = df_audio.groupby('emotion')[key_features].mean()
print(summary stats.round(2))
```

Audio Features by Emotion



Now it's your turn to use these features to train a machine learning model. Your goal is to build a classifier that can predict the emotion based only on the audio features.

Tasks:

- 1. Prepare the data by separating features (X) and labels (y).
- 2. Split the data into training and testing sets. Remember to use stratify to keep the emotion distribution even.
- 3. Scale the features using StandardScaler. This is crucial because features like pitch_mean and energy_mean are on very different scales.
- 4. Train a RandomForestClassifier on the scaled training data.
- 5. Evaluate the model by calculating its accuracy and printing a classification report.
- 6. Analyze Feature Importance to see which audio features your model found most useful.

```
# / YOUR TURN: Train and analyze an audio emotion classifier
# 1. Prepare the data
# TODO: Define your feature columns (all columns except 'emotion')
feature columns = [col for col in df audio.columns if col != 'emotion']
X audio = df audio[feature columns]
y_audio = df_audio['emotion']
# 2. Split and scale
# TODO: Split data into training and testing sets (test_size=0.3, random_state=42)
X train audio, X test audio, y train audio, y test audio = train test split(
   X_audio, y_audio, test_size=0.3, random_state=42, stratify=y_audio
# TODO: Initialize a StandardScaler and scale your training and test data
scaler = StandardScaler()
X train audio scaled = scaler.fit transform(X train audio)
X_test_audio_scaled = scaler.transform(X_test_audio)
# 3. Train classifier
# TODO: Initialize and train a RandomForestClassifier
audio classifier = RandomForestClassifier(n estimators=100, random state=42)
audio_classifier.fit(X_train_audio_scaled, y_train_audio)
# 4. Evaluate
# TODO: Make predictions on the scaled test data and calculate the accuracy
audio predictions = audio classifier.predict(X test audio scaled)
audio accuracy = accuracy score(y test audio, audio predictions)
print("  YOUR Audio Emotion Classification Results:")
print(f"Accuracy: {audio_accuracy:.3f}")
print("\nClassification Report:")
print(classification report(y test audio, audio predictions))
# 5. Feature importance analysis
# TODO: Get feature importances from the trained model
importances = audio classifier.feature importances
feature importance df = pd.DataFrame({
   'feature': feature columns,
    'importance': importances
}).sort_values('importance', ascending=False)
print("\n \ Most Important Audio Features in YOUR Model:")
print(feature importance df.head(5))
# 6. Visualize feature importance
```

```
plt.figure(figsize=(10, 6))
sns.barplot(x='importance', y='feature', data=feature_importance_df.head(5), palette='viridis')
plt.title('Top 5 Most Important Audio Features')
plt.xlabel('Importance Score')
plt.ylabel('Audio Feature')
plt.show()
```

→ YOUR Audio Emotion Classification Results: Accuracy: 1.000

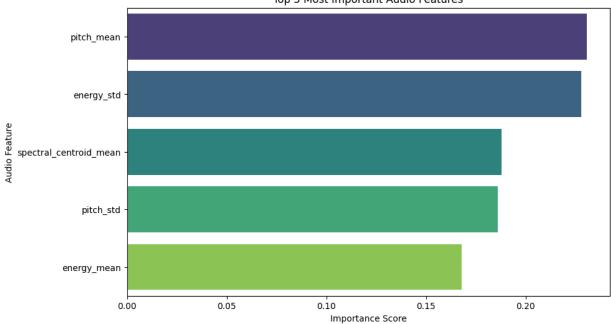
Classification Report:

	precision	recall	f1-score	support
	•			
angry	1.00	1.00	1.00	8
happy	1.00	1.00	1.00	7
neutral	1.00	1.00	1.00	8
sad	1.00	1.00	1.00	7
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

Most Important Audio Features in YOUR Model:

	i ca cui c	Impor carree
0	pitch_mean	0.230567
3	energy_std	0.227874
4	spectral_centroid_mean	0.187737
1	pitch_std	0.186095
2	energy mean	0.167727

Top 5 Most Important Audio Features



Answer based on YOUR specific results and visualizations above:

Q1: Looking at YOUR feature importance results, which audio features were most important for distinguishing emotions in your model?

Your answer based on your results: The pitch_mean were the most important for distingushing emotions.

Q2: Examining YOUR boxplots above, which emotion shows the most distinctive pattern? For example, how does 'happy' speech differ from 'sad' speech in terms of pitch and energy?

Your observation of your visualizations: "Happy" is higher on pitch_mean than "sad" and is way higher on the energy side as well.

Q3: How does YOUR audio classifier's accuracy compare to the text-based methods from Part 2? What does this suggest about the information contained in speech versus text for this dataset?

Your comparison and interpretation: The accuracy of my audio classifier performed better than my text classifier. I believe this is because there is direct association to the sounds being produced. It is harder to classify text whenever the sentiment is ambiguous.

Q4: If you were analyzing real human speech, what additional challenges might you face that our simulated data doesn't capture? (Think about background noise, different speakers, accents, etc.)

Your answer: I think challenges could be that people have different tones associated with emotions. Some people do not yell when angry, while others may yell when they are sad. This can be classified incorrectly.

Part 4: Multimodal Emotion Analysis

Combining All Previous Modules: The Ultimate Integration

This is the final and most important part of our lab. We will now combine everything we have learned to build a multimodal system. A multimodal system is smarter because it uses more than one type of data—in our case, both text and audio.

- · Module 2 (Preprocessing): We'll need clean text.
- Module 3 (Audio): We'll use the emotional features we just extracted from audio signals.
- · Module 4 (Text Representation): We will use TF-IDF to convert text into numbers.
- · Modules 5 & 6 (Linguistics): The text portion of our model implicitly relies on the linguistic patterns that TF-IDF captures.
- · Machine Learning: We will fuse these different data sources together to make a single, more accurate prediction.
- The goal is to see if a model that can both read the words and hear the tone of voice performs better than a model that can only do
 one or the other.

```
# Create multimodal dataset (text + corresponding audio)
multimodal texts = {
    'happy': [
       "I'm so excited about this opportunity!",
        "This is absolutely wonderful news!",
        "I love spending time with my family!",
        "What a beautiful day it is today!",
        "I'm thrilled to be here with you all!"
    ],
    'sad': [
        "I'm feeling really down today.",
        "This situation makes me very sad.",
       "I miss my old friends so much.",
        "Everything seems to be going wrong.",
        "I feel so lonely and isolated."
    ],
    'angry': [
```

```
"This is completely unacceptable!",
       "I'm furious about this decision!",
       "How dare you treat me this way!",
       "This makes me so angry and frustrated!",
       "I can't believe this is happening!"
   ],
    'neutral': [
       "The meeting is scheduled for tomorrow.",
       "Please submit your report by Friday.",
       "The weather forecast shows rain.",
       "I need to buy groceries later.",
       "The train arrives at 3:30 PM."
# Create multimodal dataset
multimodal data = []
print(" \( \subseteq \) Creating Multimodal Dataset (Text + Audio):")
for emotion, texts in multimodal_texts.items():
   print(f"\n [ {emotion.capitalize()} examples:")
   for i, text in enumerate(texts):
       # Generate corresponding audio using the function from Part 3
       audio, sr = create_emotional_audio(emotion, duration=2.0)
       # Extract audio features using the function from Part 3
       audio features = extract emotional features(audio, sr)
       # Store the text, the audio features, and the emotion label together
       multimodal_data.append({
           'text': text,
           'audio features': audio features,
           'emotion': emotion
       print(f" Text: {text}")
       if i == 0: # Play the first audio sample for demonstration
           print(f" Listen to its corresponding audio:")
           display(Audio(audio, rate=sr))
print("Each sample now contains both text and its corresponding emotional audio features.")
```

```
→ Treating Multimodal Dataset (Text + Audio):

☐ Happy examples:

      Text: I'm so excited about this opportunity!
      Listen to its corresponding audio:
          0:00 / 0:02
      Text: This is absolutely wonderful news!
      Text: I love spending time with my family!
      Text: What a beautiful day it is today!
      Text: I'm thrilled to be here with you all!

√ Sad examples:
      Text: I'm feeling really down today.
      Listen to its corresponding audio:
          0:00 / 0:02
      Text: This situation makes me very sad.
      Text: I miss my old friends so much.
      Text: Everything seems to be going wrong.
      Text: I feel so lonely and isolated.
    Angry examples:
      Text: This is completely unacceptable!
      Listen to its corresponding audio:
          0:02 / 0:02
      Text: I'm furious about this decision!
      Text: How dare you treat me this way!
      Text: This makes me so angry and frustrated!
      Text: I can't believe this is happening!

■ Neutral examples:
      Text: The meeting is scheduled for tomorrow.
      Listen to its corresponding audio:
          0:00 / 0:02
      Text: Please submit your report by Friday.
      Text: The weather forecast shows rain.
      Text: I need to buy groceries later.
      Text: The train arrives at 3:30 PM.
    Multimodal Dataset Created: 20 samples
    Each sample now contains both text and its corresponding emotional audio features.
```

4.1 Fusing Text and Audio Features

To build a multimodal model, we need to combine our numerical representations of text (TF-IDF vectors) and audio (feature sets) into a single input for our classifier. This process is called feature fusion. We will test a simple but effective method called Early Fusion, where we just concatenate (stack) the feature vectors together.

```
# Main Analysis: Build and Compare Multimodal Emotion Detection Systems
# 1. Extract features from both modalities
multimodal_texts_list = [sample['text'] for sample in multimodal_data]
multimodal_emotions = [sample['emotion'] for sample in multimodal_data]
```

```
# Text features (Module 4 concepts)
multimodal_tfidf = TfidfVectorizer(max_features=50, stop_words='english')
text_features = multimodal_tfidf.fit_transform(multimodal_texts_list).toarray()
# Audio features (Module 3 concepts)
# We need to extract the audio features from our list of dictionaries
audio features df = pd.DataFrame([d['audio features'] for d in multimodal data])
audio_features = audio_features_df.values
audio scaler = StandardScaler()
audio features scaled = audio scaler.fit transform(audio features)
# 2. Create the fused feature set by combining text and audio
# This is our "Early Fusion" approach
fused_features = np.concatenate([text_features, audio_features_scaled], axis=1)
print(f"Shape of Text Features: {text features.shape}")
print(f"Shape of Audio Features: {audio features scaled.shape}")
print(f"Shape of Fused Features: {fused_features.shape}")
# 3. Train and evaluate different approaches
approaches = {
    'Text Only': text_features,
    'Audio Only': audio_features_scaled,
    'Multimodal (Fused)': fused_features
results = {}
print("\n \sqrt{n \sqrt{n}} Comparing Model Performance:")
print("=" * 50)
for approach name, features in approaches.items():
   # Split data
    X_train, X_test, y_train, y_test = train_test_split(
        features, multimodal_emotions, test_size=0.3, random_state=42, stratify=multimodal_emotions
    # Train a Logistic Regression classifier
    classifier = LogisticRegression(random state=42, max iter=1000)
    classifier.fit(X_train, y_train)
    # Evaluate
    predictions = classifier.predict(X test)
    accuracy = accuracy score(y test, predictions)
    results[approach_name] = accuracy
    print(f"{approach_name:20}: Accuracy = {accuracy:.3f}")
# 4. Analyze the improvement from fusion
best single modality acc = max(results['Text Only'], results['Audio Only'])
multimodal acc = results['Multimodal (Fused)']
improvement = multimodal_acc - best_single_modality_acc
print("\n Analysis of Multimodal Improvement:")
print(f"Best Single-Modality Accuracy: {best single modality acc:.3f}")
print(f"Multimodal (Fused) Accuracy: {multimodal_acc:.3f}")
print(f"Improvement from Fusion:
                                      {improvement:+.3f}")
if improvement > 0.01: # Check for a meaningful improvement
    print("\n✓ Conclusion: Multimodal fusion provided a clear improvement over using just text or audio alone!")
else:
    print("\n ▲ Conclusion: Multimodal fusion did not provide a significant improvement in this case.")
```

o YOUR TURN - Final Comprehensive Reflection

This is the final part of the lab. Based on everything you have built and the results you see above, answer the following questions.

Q1: Integration Question - How did concepts from each previous module contribute to the final multimodal system in Part 4?

- · Module 2 (Preprocessing): Normalization, noise reduction, tokenization, lowercasing, and data alignment all contributed.
- · Module 3 (Audio): A great representation of the tone, pitch, energy gave the system the ability to deliver the accuracy results needed.
- Module 4 (Text Representation): The ability to understand context led to a greater grasps of intent, sentiment, and fusion of audio with text

Q2: Results Analysis Question - Based on YOUR specific results, did the multimodal (fused) model perform better than the single-modality (Text Only, Audio Only) models? Why do you think this happened?

Your analysis of your results: The accuracy of text only was .167, while audio performed at a 1.0. The multimodel performed at a 1.0 which was not an improvement compared to audio. I believe it performed the same due to the audio being non-human. The text was based on human interaction and the accuracy was way lower.

Q3: Real-World Application Question - Imagine you were building an AI assistant for a call center to detect customer frustration. Which model would you choose (Text Only, Audio Only, or Multimodal)? Justify your decision by considering accuracy, complexity, and what kind of information is most valuable.

Your business decision and reasoning: I would use a multimodel due to human vocals and language being very comprehensive. I believe that the audio alone would lead to misunderstandings, while text alone could lead to lack of context and intent. Both together have a better shot of success. I also have to consider customers that will not be able to speak over the phone due to disabilities. Rather there can be numbers pressed associated with emotional review. Such as 1 being happy, 2 being upset, 3 being needing help, etc...

Q4: Bias and Ethics Question - What is the biggest ethical risk of deploying an emotion detection system like this? Describe one potential type of bias (e.g., cultural, gender, age) and explain how it might cause problems.

Your answer on ethics and bias: There could be ethical concerns regarding this detection system due to the variety of human communication. There are language barriers, differences in speech/lingo, and older individuals having a more mature voice than someone younger. There are a lot of factors that could complicate this project. The largest ethical risk is not being universal.

Q5: Personal Experience Question - What was the most challenging part of this lab for YOU personally, and what was the most surprising or interesting discovery you made?

Your personal experience and discoveries: The most challenging part was finding real world examples of sentiment ambiguity. I struggled with this part, but once I was able to translate it into human actions, it became easier. The most surprising thing I learned was that there can be audio generated and analyzed without it being a recorded audio!

Lab Conclusion

What You've Accomplished

Congratulations! You've successfully:

- 1. Connected all previous modules to build a comprehensive emotion analysis system
- 2. **Built rule-based sentiment analyzers** using VADER and TextBlob
- 3. Created machine learning classifiers using TF-IDF and scikit-learn
- 4. **☑ Developed speech emotion detection** using audio feature extraction
- 5. Implemented multimodal fusion combining text and audio analysis
- 6. ✓ Analyzed bias and ethical considerations in emotion detection systems

Key Insights from Your Journey

- · Integration is powerful: Combining concepts from all modules creates more robust systems
- Different approaches have different strengths: Rule-based, ML, and multimodal each excel in different scenarios
- · Context matters: The same words can convey different emotions depending on how they're spoken
- · Bias is real: Emotion detection systems can perpetuate societal biases and must be carefully evaluated
- · Practical considerations: Real-world deployment requires balancing accuracy, interpretability, and computational cost

Looking Forward

This lab represents the culmination of your foundational NLP learning. In upcoming modules, you'll learn about:

- Neural networks for more sophisticated emotion analysis
- Deep learning architectures that can learn complex patterns
- · Transformer models that understand context even better
- Large language models that can perform emotion analysis with minimal training

Final Thought